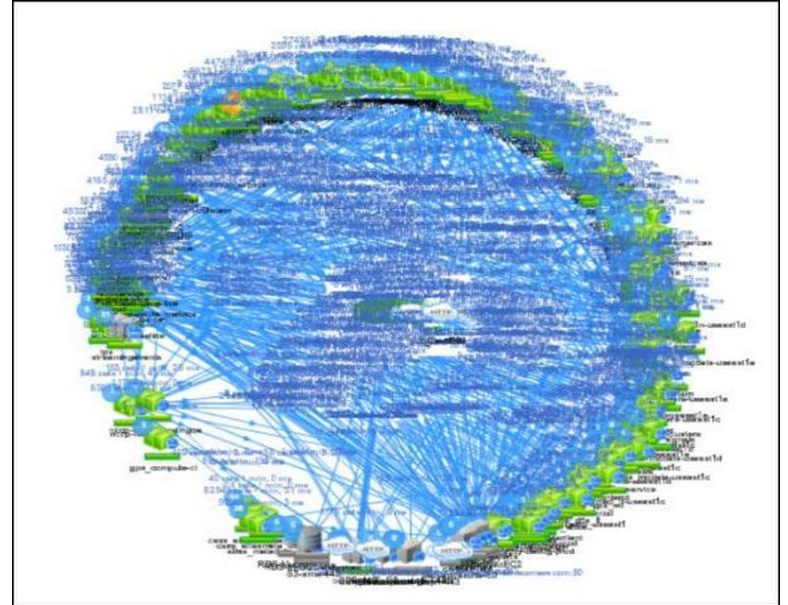

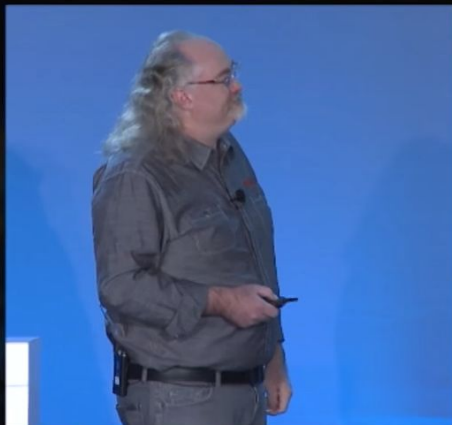



# Microservices








## Netflix architecture



© 2019 Amazon.com, Inc. and its affiliates. All rights reserved. May not be copied, modified, or distributed in whole or in part without the express consent of Amazon.com, Inc.


# # whoami

- Martin Štefanko
- Software engineer, Red Hat
- MicroProfile committer
- Microservices enthusiast
-  @xstefank

# What we are going to go through?

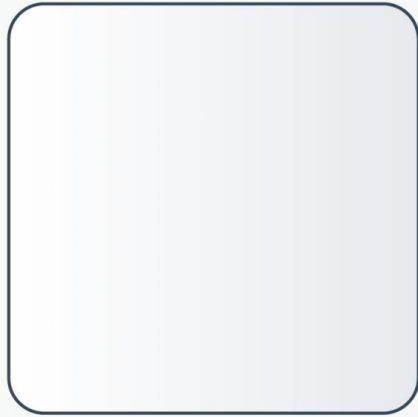
- 1st lecture – Microservices introduction
- 2nd lecture – Introduction to Quarkus, HTTP, REST, Docker/Podman, and Kubernetes/OpenShift
- 3rd lecture – MicroProfile, Jakarta EE
- 4th lecture – MicroProfile continued, Panache, reactive programming, (Kafka)

# Microservices



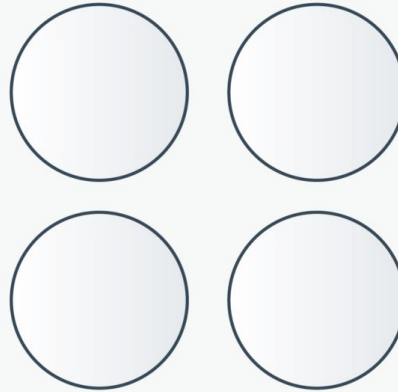


# Monolithic vs. SOA vs. Microservices




**Monolithic**

**Single Unit**



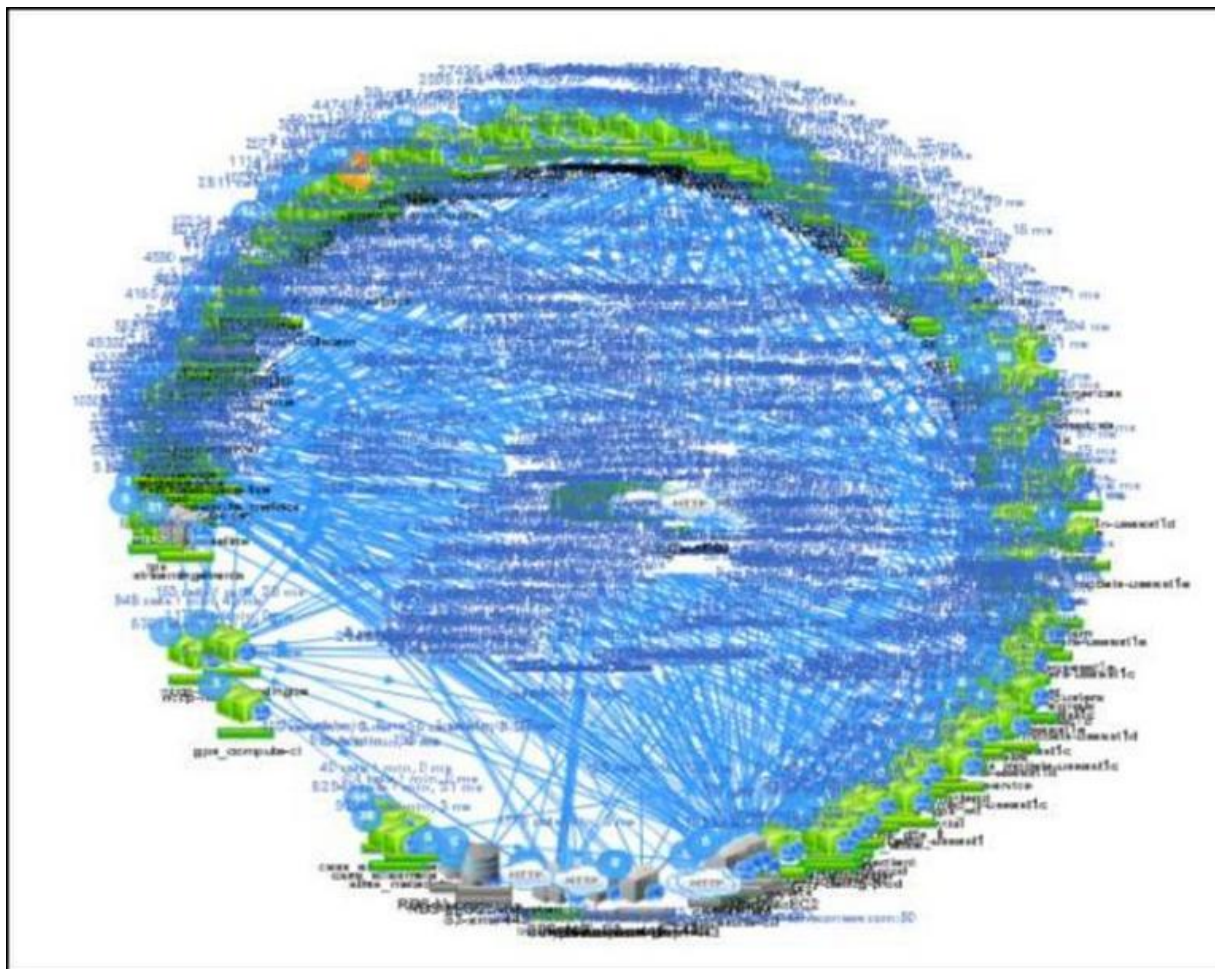
**SOA**

**Coarse-grained**

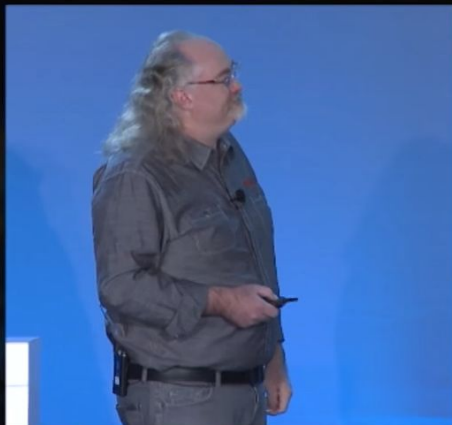


**Microservices**

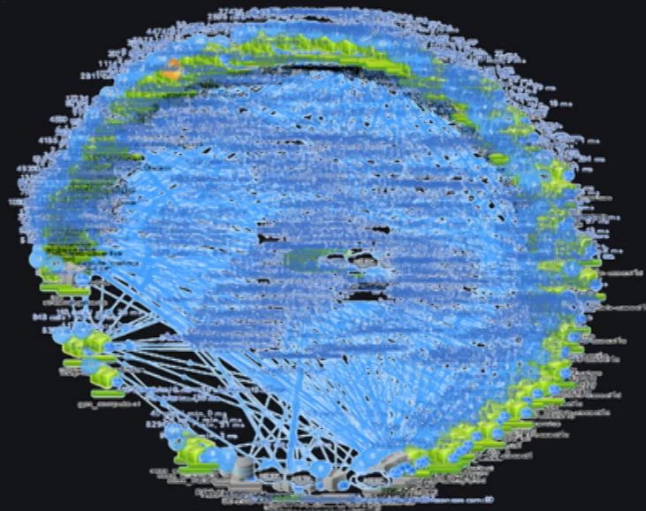
**Fine-grained**



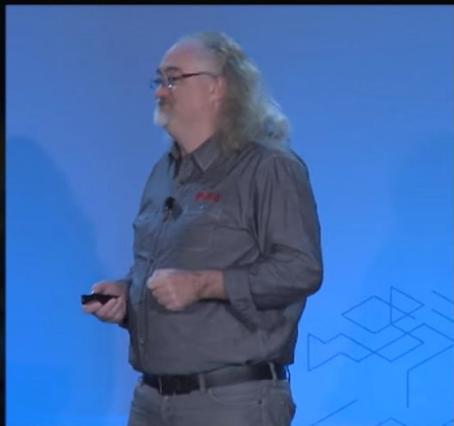




## Netflix architecture



© 2019 Amazon.com, Inc. and its affiliates. All rights reserved. May not be copied, modified, or distributed in whole or in part without the express consent of Amazon.com, Inc.



## Netflix ecosystem

100s of microservices

1000s of daily production changes

10,000s of instances

100,000s of customer interactions per minute

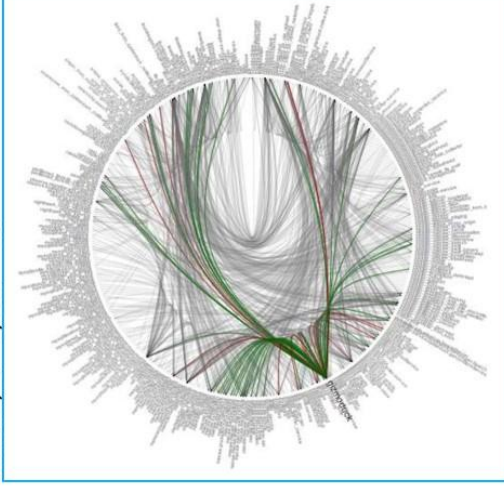
1,000,000s of customers

1,000,000,000s of metrics

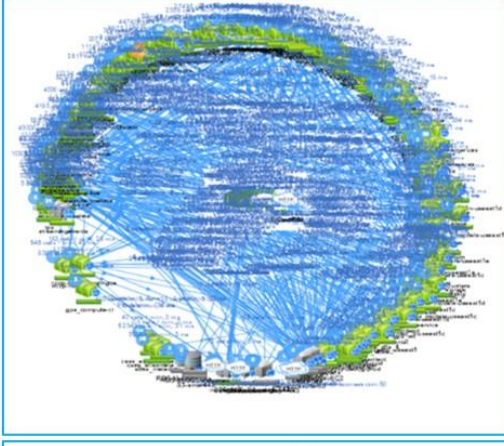
10,000,000,000 hours of streamed

**10s of operations engineers**

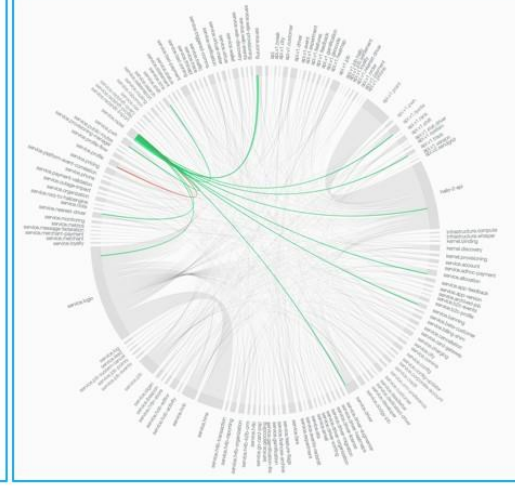
Twitter, 2013, 500+ ms



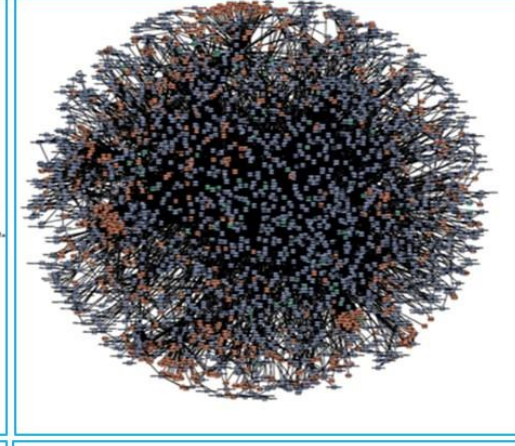
Netfix, 500+ ms



Hailo, 450+ ms



Amazon, 2009



# Monolith

# Monolithic applications

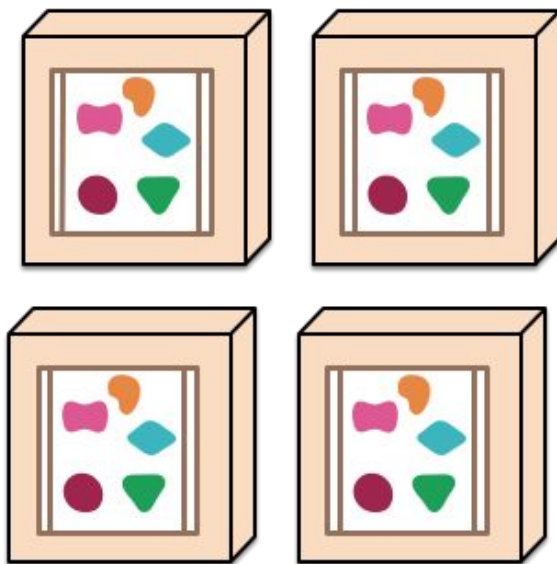
- Common development model (past 20+ years)
- Single or small amount of deployments
- Application server
- Single process
- Components tightly coupled



*A monolithic application puts all its functionality into a single process...*



*... and scales by replicating the monolith on multiple servers*



# Monolith - advantages

- Development model == application requirements
  - Traditionally CRUD or MVC
  - Presentation - Business layer - Database
- Single or small number of archives/deployments
- Easy horizontal scaling

# Monolith - problems

- Adding new functionality
- General maintenance (bug fixing, CVEs)
- Every single change means requires rebuild and redeploy of the whole application
- Replicated server instances take more resources
- Slow startup times

# Monolith - problems

- Often extremely large code bases
  - Hard to understand / maintain
  - Long learning curves
  - Experts in particular system / part of the system
  - Often a commitment to a particular technology
    - Or even a specific version

# Microservices



# Architectural pattern

- System as a collection of small, isolated services
- Each service
  - Owns its data
  - Is independently isolated
  - Is scalable
  - Is resilient
  - Is self-maintained
- Evolution of SOA
- Intuitive approach


# Architectural pattern

- Each service represents the separated and independent part of the system
- Interaction with other microservices is allowed only through predefined communication interfaces (API)


*A monolithic application puts all its functionality into a single process...*




*... and scales by replicating the monolith on multiple servers*



*A microservices architecture puts each element of functionality into a separate service...*



*... and scales by distributing these services across servers, replicating as needed.*



# Organization structure

- Teams segregation according to the architecture
- Each team is responsible for one or a small set of services
- Structured according to the business goals
  - Teams should be self-maintained
  - Devs, Testers, Front-end, ...
- One team should not have access/knowledge of different services/teams

# Isolation

- Ownership of resources
- Data requests to different services prohibited only through the API
  - Control the access, computation requirements
- Service must act as an external component
- Loosely coupled services
- Virtual addresses
- Scaling, load-balancing
- Resiliency



# Isolation

- Underlying technology may differ
  - Different languages
  - Different runtimes
  - Different frameworks
  - Different versions

# Law of Demeter

- Microservice typically must communicate with other microservices to provide its functionality
- Principle of least knowledge
  - Each unit should have only limited knowledge about other units: only units "closely" related to the current unit
  - Each unit should only talk to its friends; don't talk to strangers

# Single responsibility principle

- “*micro*” doesn’t necessarily mean small
- Micro == scope of the service responsibility
- SRP – a class (microservice) should have only one reason to change
- Unix philosophy – Make each program (microservice) do one thing well

# API

- Application programming interface
- Technology-agnostic
- Remote procedure calls
  - HTTP & REST
  - Apache Kafka
  - gRPC
  - AMQP, MQTT, JMS
  - ....

```
1  swagger: "2.0"
2  info:
3    version: "1.0.0"
4    title: Hello World App
5    host: api.hello-world-example.com
6    basePath: /
7  schemes:
8    - https
9  paths:
10   /hello:
11     get:
12       description: Returns 'Hello' to the caller
13       parameters:
14         - name: name
15           in: query
16           description: The name of the person to whom to say hello
17           required: false
18           type: string
19       responses:
20         200:
21           description: OK
22         default:
23           description: Error
24
```




# Microservices vs SOA

# Monolithic vs. SOA vs. Microservices




**Monolithic**

**Single Unit**



**SOA**

**Coarse-grained**



**Microservices**

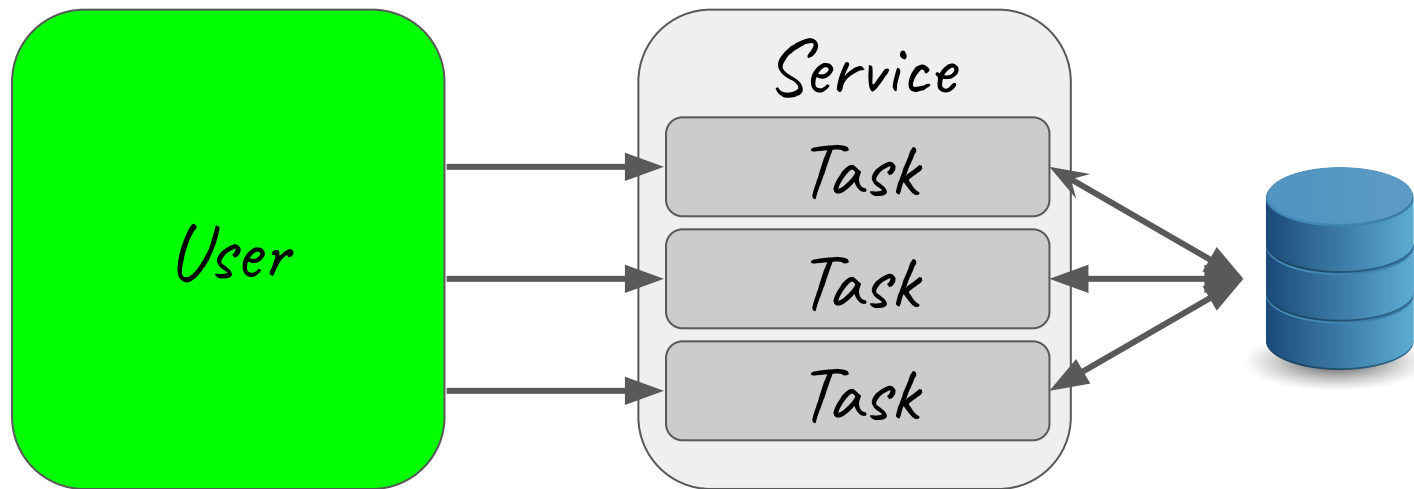
**Fine-grained**

# Microservices vs SOA

- Both are about services...
- but the service characteristics differ

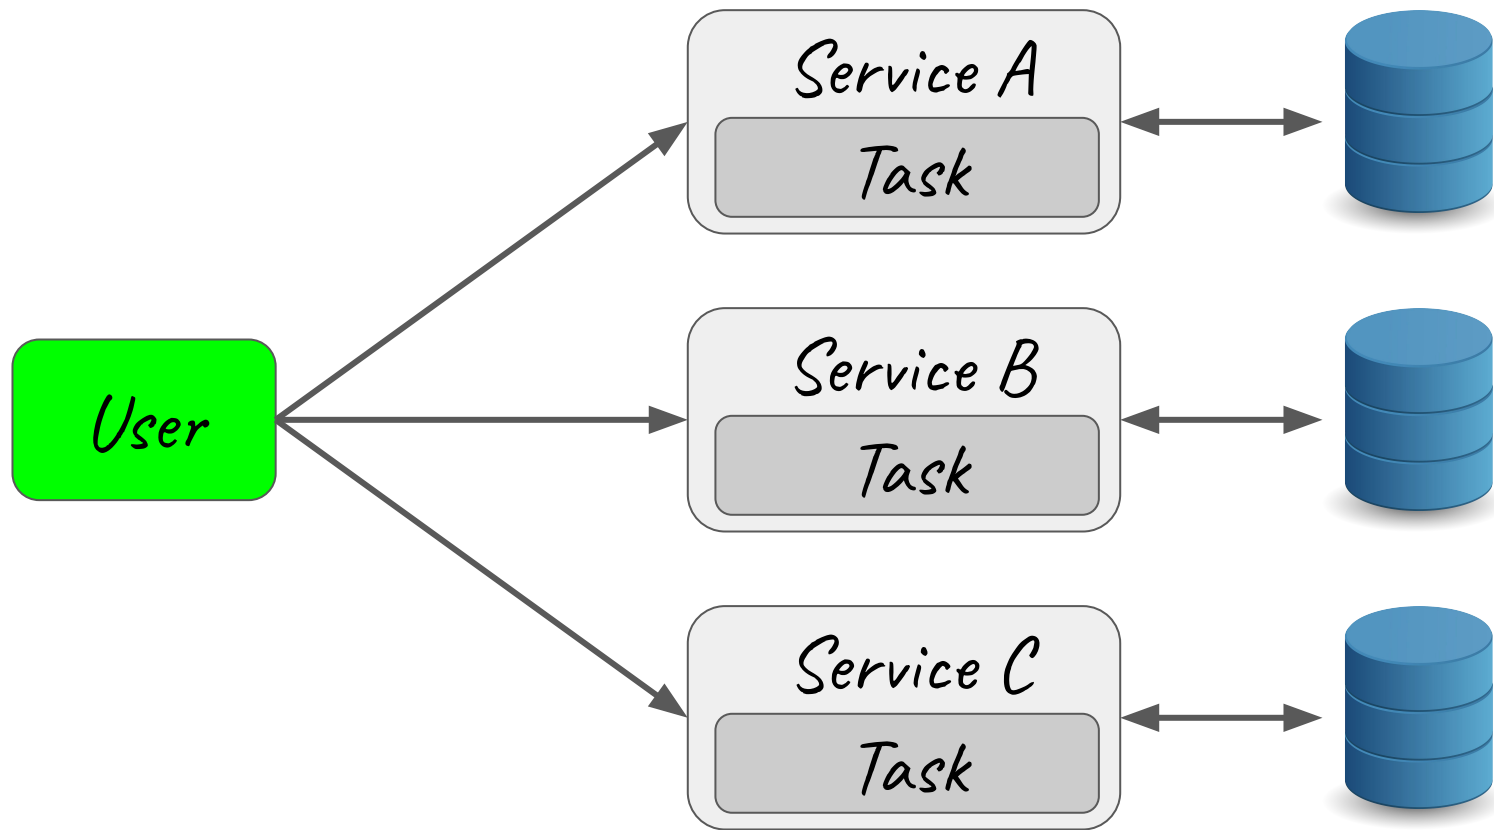
# SOA

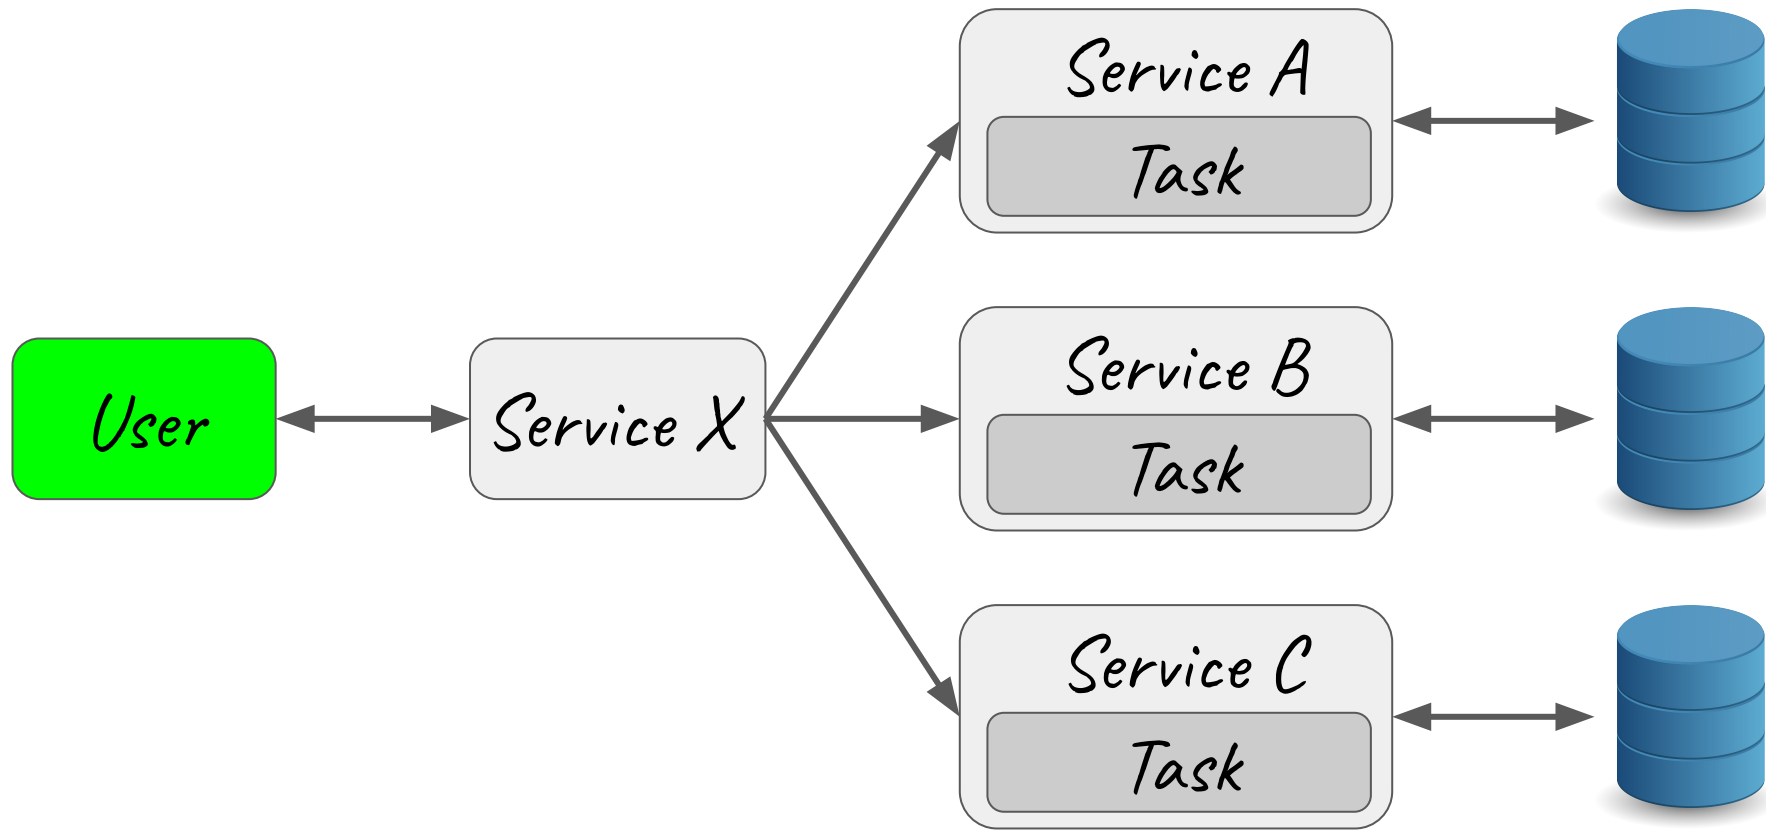
- Smaller number of services
- Each service is responsible for a particular group of functions/tasks that together provide a system functionality
- User requests are typically processed by one or in a small number of services
  - User can interact with the same service in several requests/steps



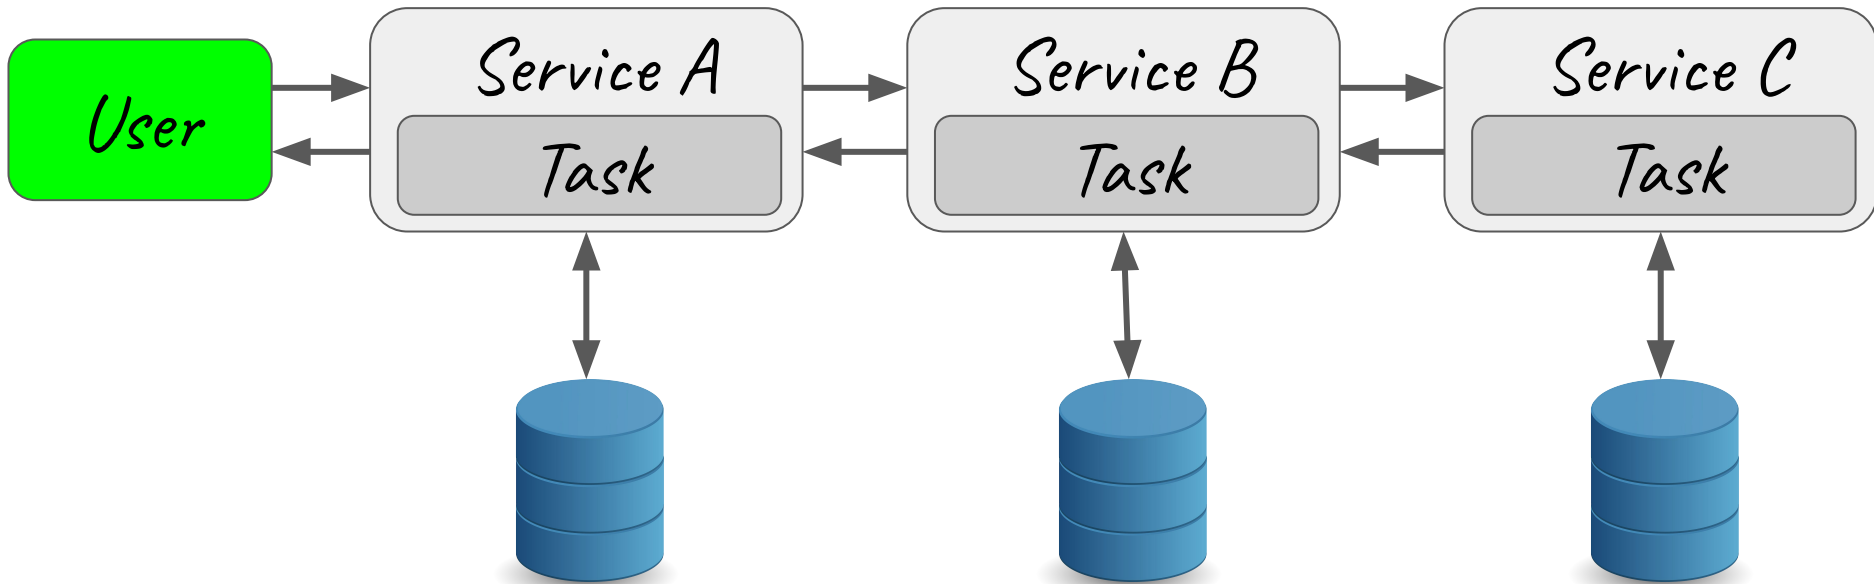
# Microservices

- Higher number of services
- Each service is responsible for one particular function/task
- Users typically interacts with several services (not necessarily directly)
- The system provides its functions as an collaboration/interaction between microservices











# Differences side-by-side

# Service organization

*SOA*

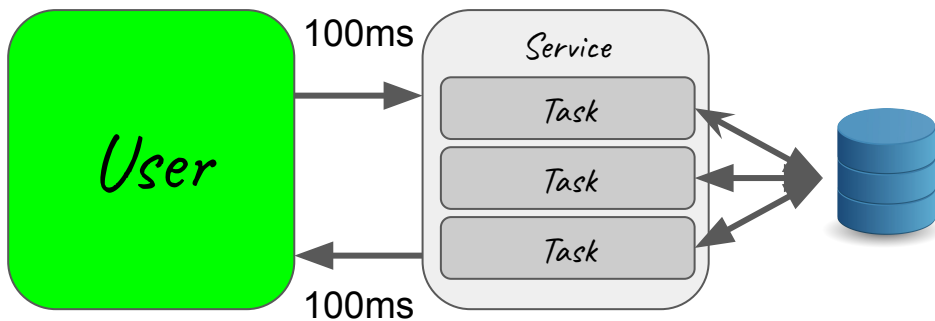


*Microservices*

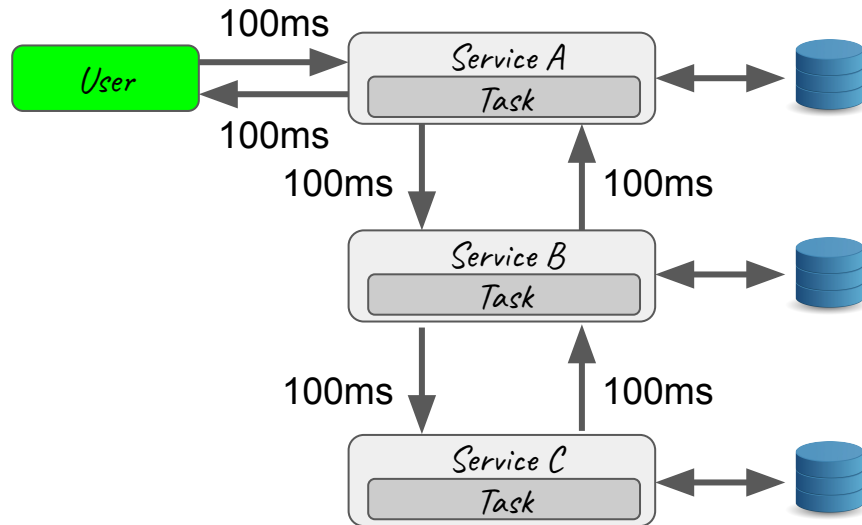


# Task granularity

*SOA*



*Microservices*



# Component sharing

*SOA*

*Microservices*

*Customer  
management*

*Warehouse  
management*

*Shipping  
management*

*Order service*

*Customer  
management*

*Warehouse  
management*

*Shipping  
management*


*Customer orders*

*Warehouse  
orders*


*Shipping orders*

# Message exchange

*SOA*



*Microservices*



## *SOA*

- Share as much as possible
- Importance on the business functionality reuse
- Common governance and standards

## *Microservices*

- Share as little as possible
- Importance on bounded contexts
- Focus on people, collaboration, freedom of options

## *SOA*

- Enterprise service bus (ESB)
- Standardized messaging protocols
- Multi-threaded

## *Microservices*

- Direct communication (service mesh)
- Lightweight protocols (HTTP/REST, Kafka)
- Single-threaded (event loop)



## *SOA*

- Maximize service reusability
- Traditional technologies (Relation DBs)

## *Microservices*

- Focus on service decoupling
- Modern technologies, faster adoption (NoSQL)

## *SOA*

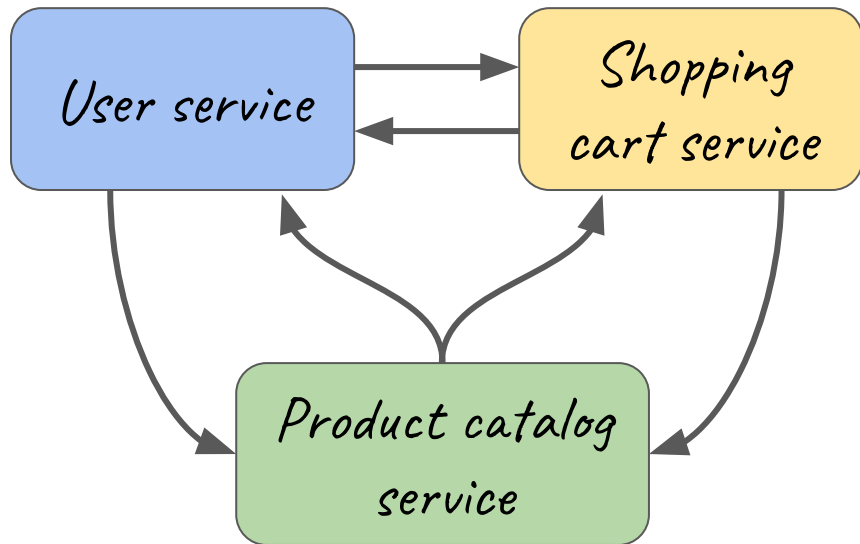
- Any change still resembles problems similar to monoliths
- Traditional development models

## *Microservices*

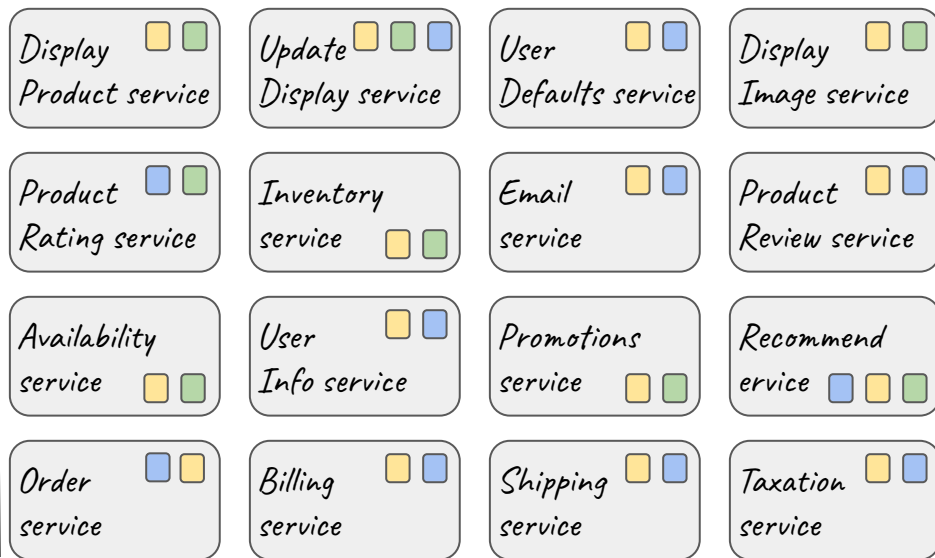
- Any change is only local to a particular service (or add a new service)
- Strong focus on **DevOps** and **CI/CD**

# Example

SOA



Microservices



# Technologies

# Docker (containers)

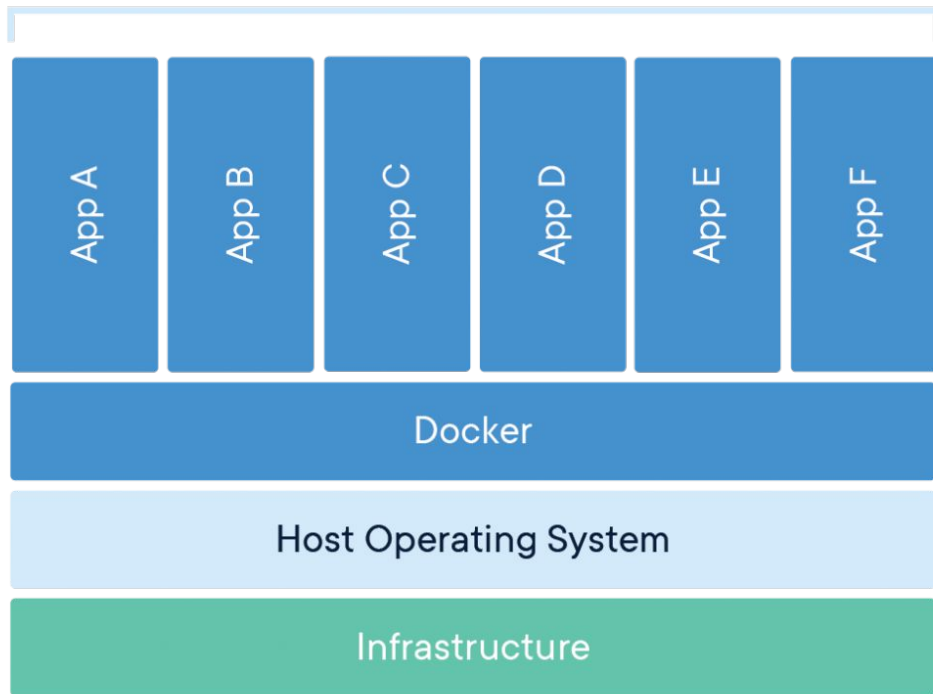


- **Container** – a standardized unit of software
- Packages code and its dependencies, runtime, system tools, system libraries, settings
- Users build **Docker images** – lightweight, standalone, executable package of software
- Images can be run anywhere where Docker is installed
- **Docker hub** ([hub.docker.com](https://hub.docker.com))

# Docker containers

- Container is runtime representation of the image
- Containers run on Docker Engine
- It doesn't matter on which platform (Linux, Mac, Windows) you run
- Containers isolate software from its environment
- Uniform behavior everywhere

## Containerized Applications



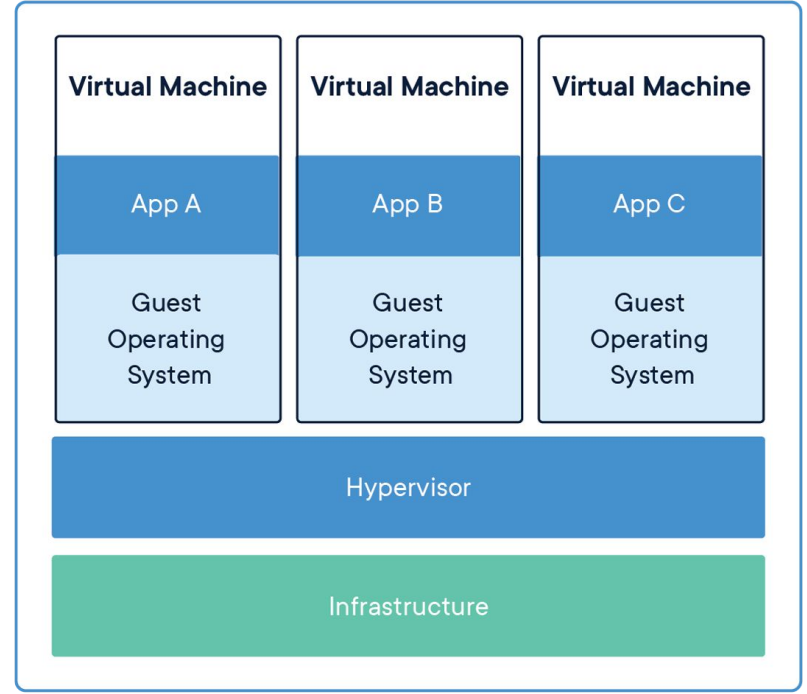
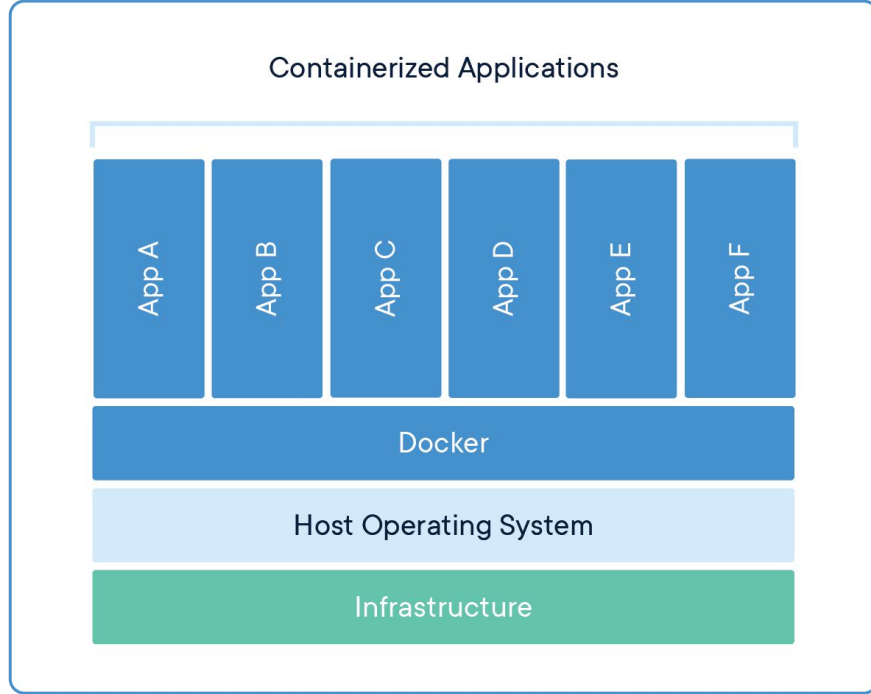
# Docker Engine

- **Standard** – Docker created the industry standard for containers, so they could be portable anywhere
- **Lightweight** – Containers share the machine's OS system kernel and therefore do not require an OS per application
- **Secure:** Applications are safer in containers and Docker provides the strongest default isolation capabilities in the industry



# Containers vs Virtual Machines

Containers and virtual machines have similar resource isolation and allocation benefits, but function differently because **containers virtualize the operating system instead of hardware**. Containers are more portable and efficient.



# Docker - standardization

- Docker launched in 2013 - revolution in application development
- In June 2015, Docker donated the container image specification and runtime code now known as runc, to the **Open Container Initiative (OCI)**
- Other alternatives - Podman, Buildah

# Dockerfile

```
FROM registry.access.redhat.com/ubi8/ubi-minimal
WORKDIR /work/
COPY target/*-runner /work/application
RUN chmod 775 /work
EXPOSE 8080
CMD [ "./application", "-Dquarkus.http.host=0.0.0.0" ]
```

```
$ docker help
```

```
Usage:  docker COMMAND
```

```
Management Commands:
```

container	Manage containers
image	Manage images

```
Commands:
```

attach	Attach to a running container
build	Build an image from a Dockerfile
create	Create a new container
pull	Pull an image or a repository from a registry
push	Push an image or a repository to a registry
run	Run a command in a new container

# Kubernetes



- Container orchestration
- an open-source system for automating deployment, scaling, and management of containerized applications
- Groups containers to logical units
- Easy administration
- De facto standard for cloud deployments

Overview

Cluster

- Cluster Roles
- Namespaces
- Nodes
- Persistent Volumes
- Storage Classes

Namespace

All namespaces

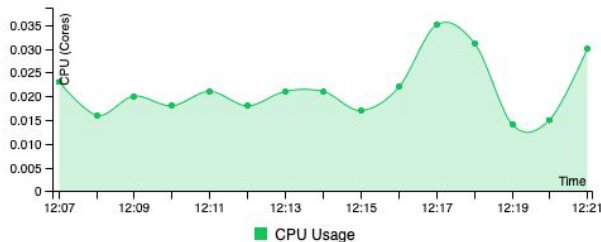
Overview

Workloads

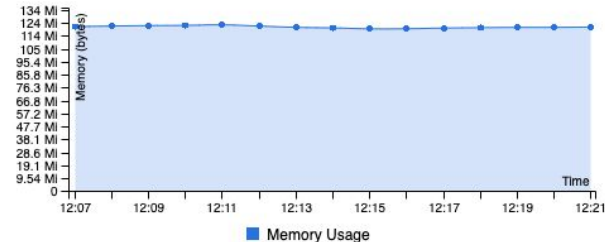
- Cron Jobs
- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets
- Replication Controllers
- Stateful Sets

Workloads

CPU Usage



Memory Usage



Workload Status



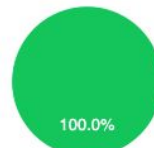
Daemon Sets



Deployments



Pods



Replica Sets

# Kubernetes objects

- **Pod** – basic executions unit
  - Process running in the cluster
  - One or multiple containers
  - Replaceable unit, can be restarted anytime (health checks)
- **Service** – exposure of application (pods) as a network service
  - Abstraction of the access to pods



# Kubernetes objects

- **Volume** – storage shared between containers in the pod
- **Deployment** – declarative updates for pods
  - User describes the desired state
  - Deployment controller (dc) changes the actual state to the desired state at controlled rate
  - New state of the pods, rollbacks, scaling,...

```
$ kubectl help
```

#### Basic Commands (Beginner):

create	Create a resource from a file or from stdin.
expose	Take a replication controller, service, deployment or pod and expose it as a new

#### Kubernetes Service

run	Run a particular image on the cluster
set	Set specific features on objects

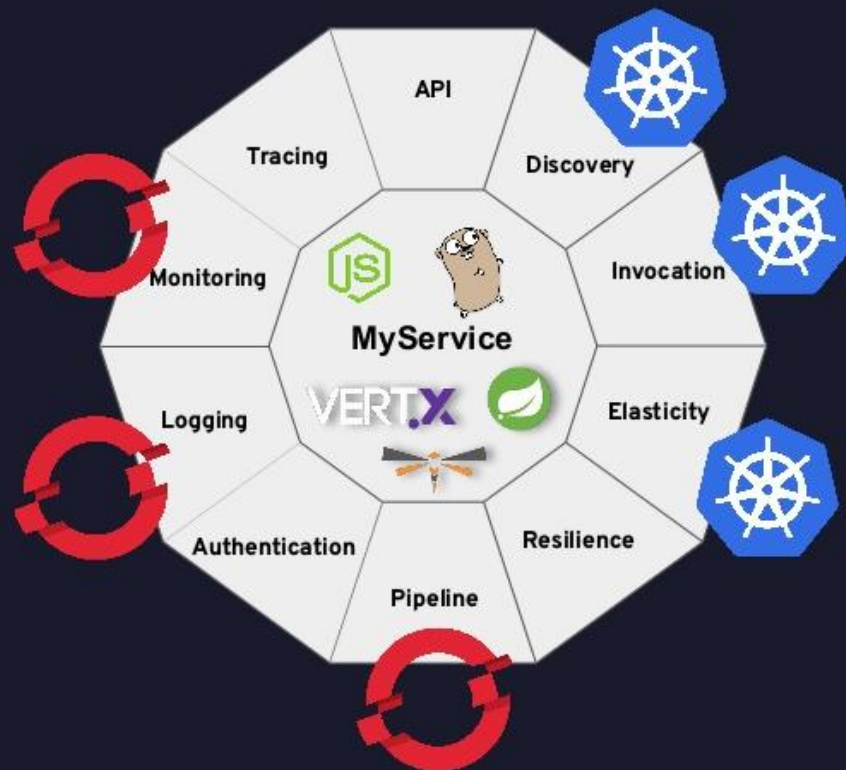
#### Deploy Commands:

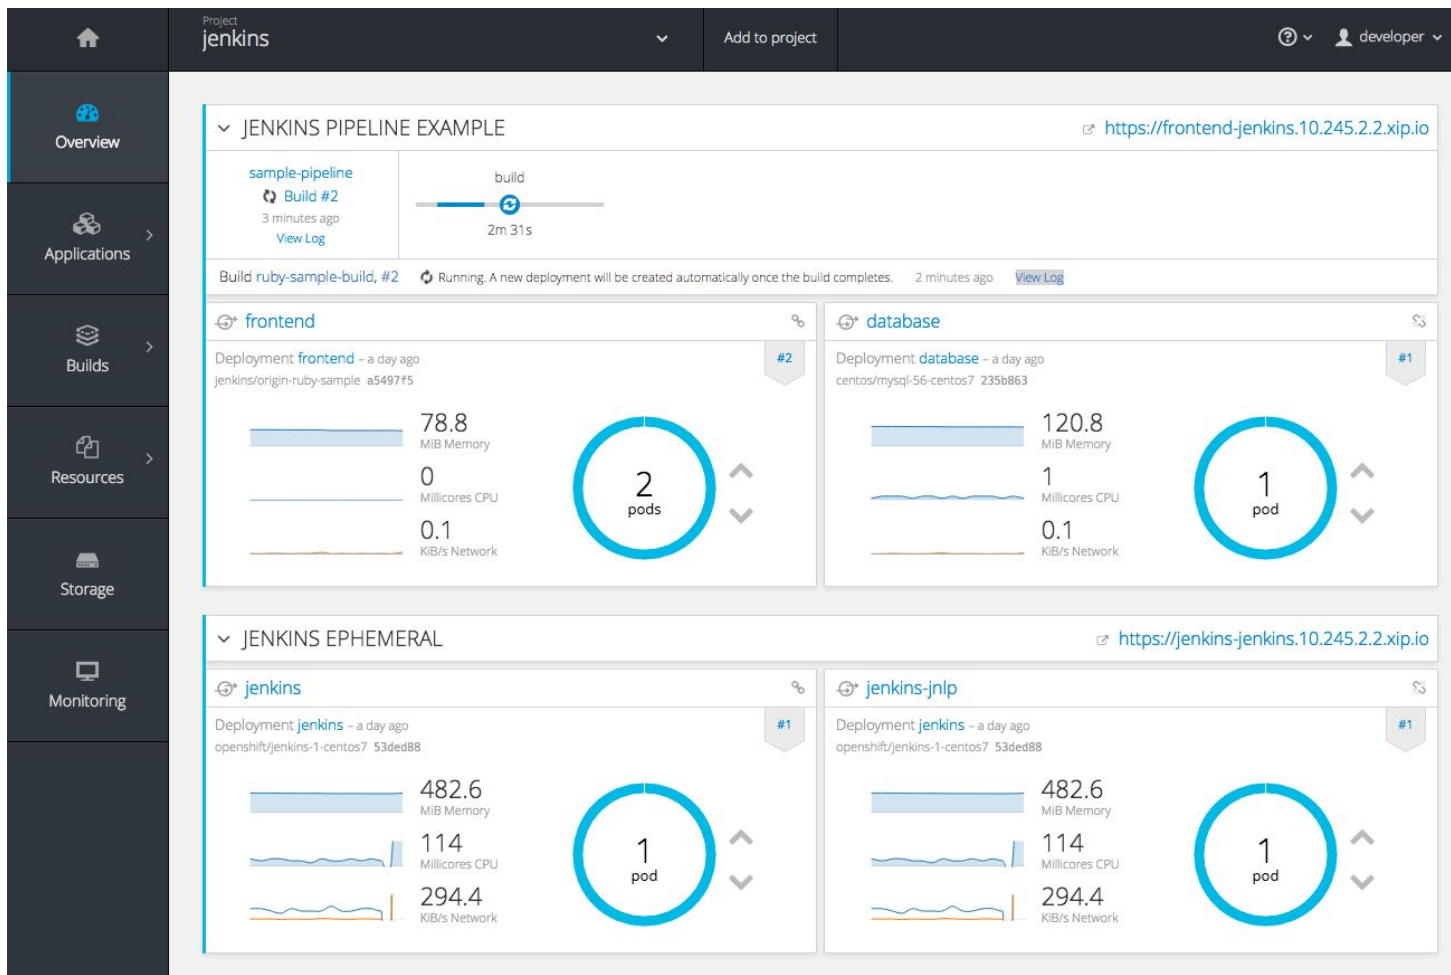
rollout	Manage the rollout of a resource
scale	Set a new size for a Deployment, ReplicaSet, Replication Controller, or Job
autoscale	Auto-scale a Deployment, ReplicaSet, or ReplicationController

# OpenShift

- Fork of Kubernetes developed and maintained at Red Hat
- Commercial product with support
- Automated installation, upgrades, and lifecycle management throughout the container stack

# Microservices'ilities + OpenShift





```
$ oc help
```

#### Basic Commands:

types	An introduction to concepts and types
new-project	Request a new project
new-app	Create a new application
status	Show an overview of the current project
project	Switch to another project
projects	Display existing projects
explain	Documentation of resources
cluster	Start and stop OpenShift cluster

#### Build and Deploy Commands:

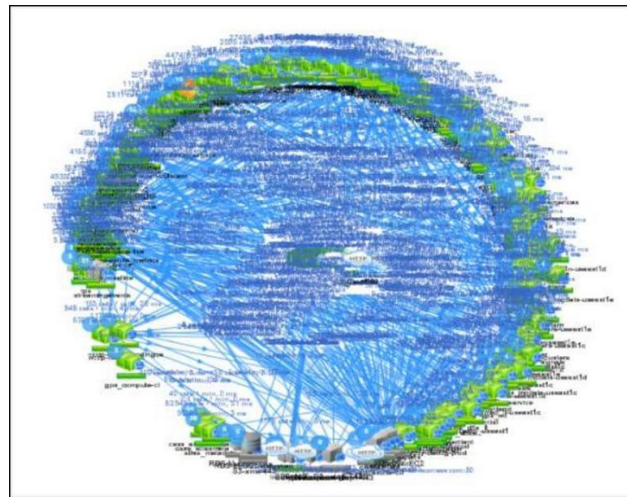
new-build	Create a new build configuration
start-build	Start a new build

#### Troubleshooting and Debugging Commands:

logs	Print the logs for a resource
------	-------------------------------

# Istio - service mesh

- **Service mesh** – the network of microservices that make up the application and the interactions between them

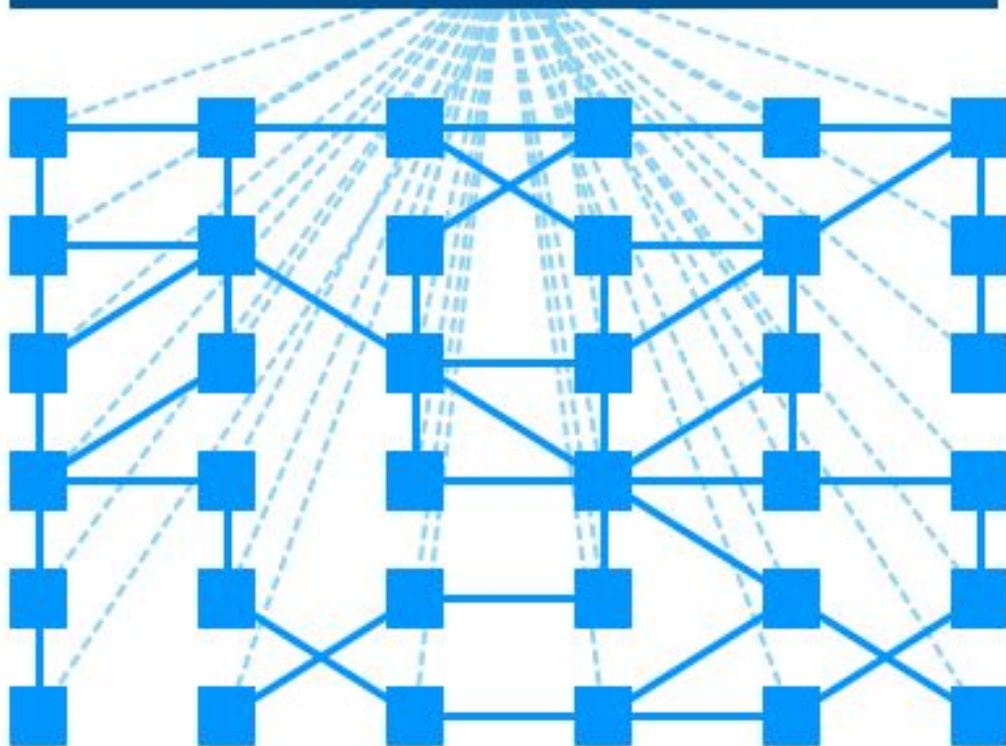


# Istio – service mesh

- As a service mesh grows in size and complexity, it can become harder to understand and manage
- requirements include discovery, load balancing, failure recovery, metrics, and monitoring
- operational requirements, like A/B testing, canary rollouts, rate limiting, access control, and end-to-end authentication

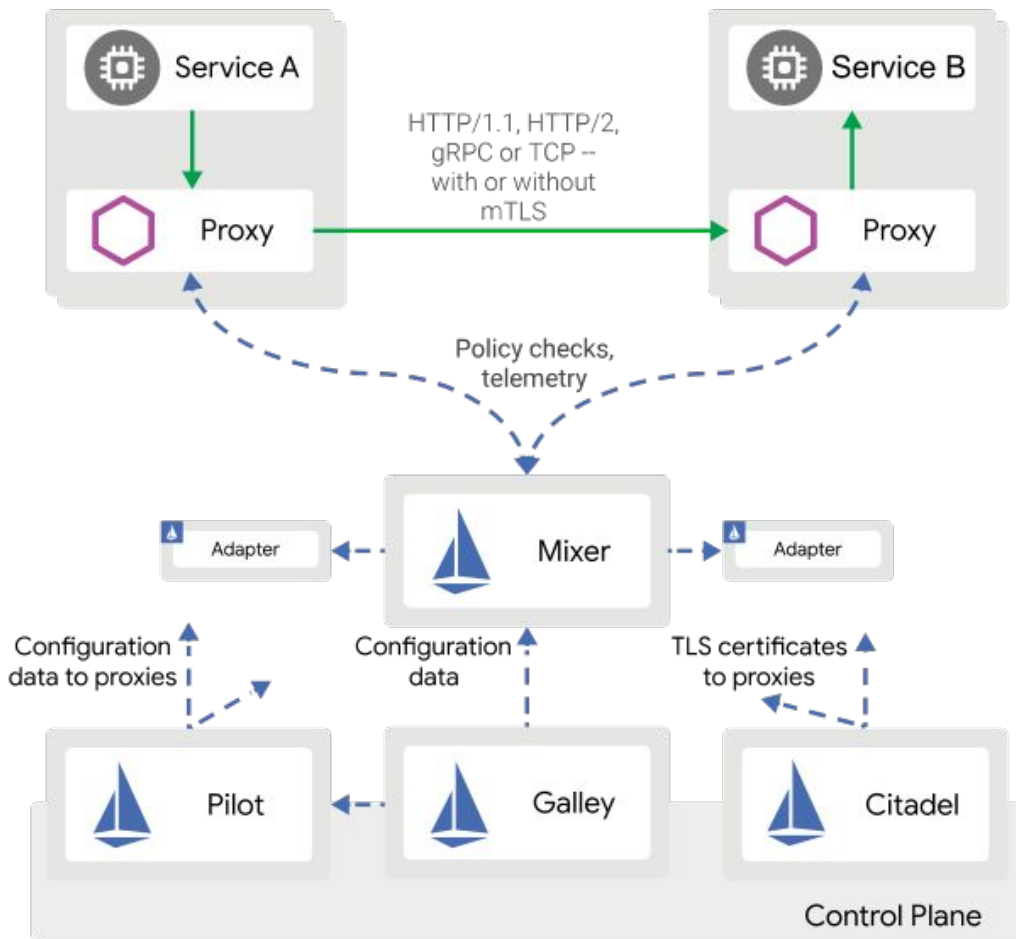


## Service Mesh's Control Plane

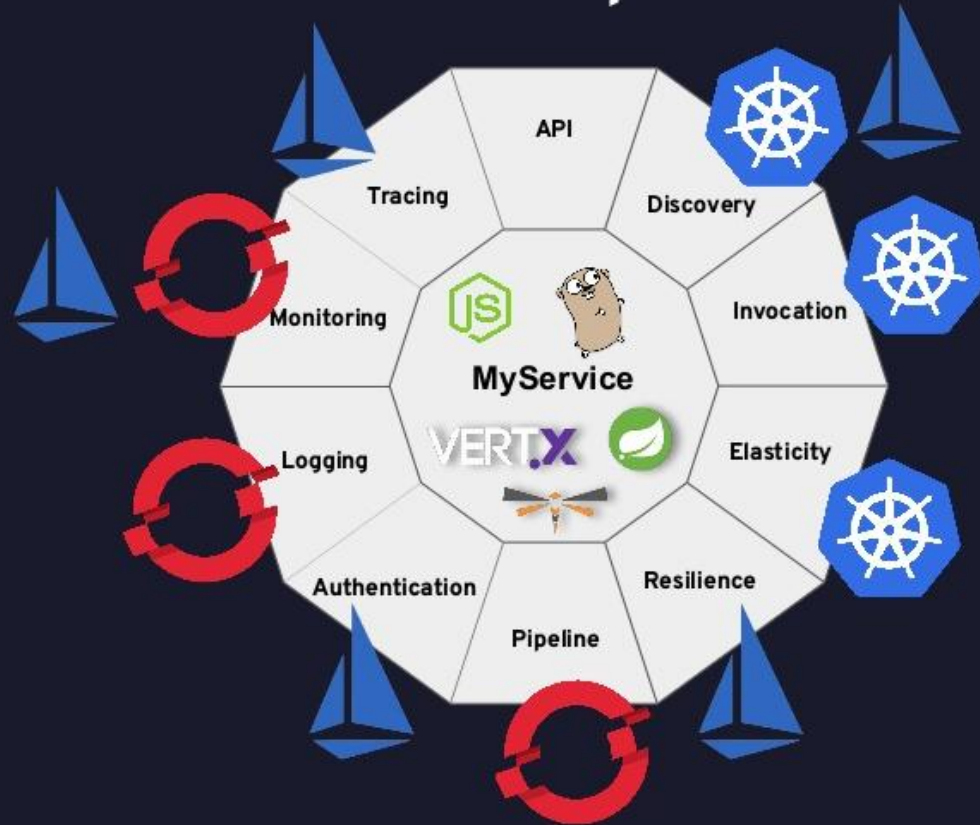


# Istio – Envoy

- Sidecar container
- Deployed in the same pod as the application container
- All network traffic goes through the Envoy proxy

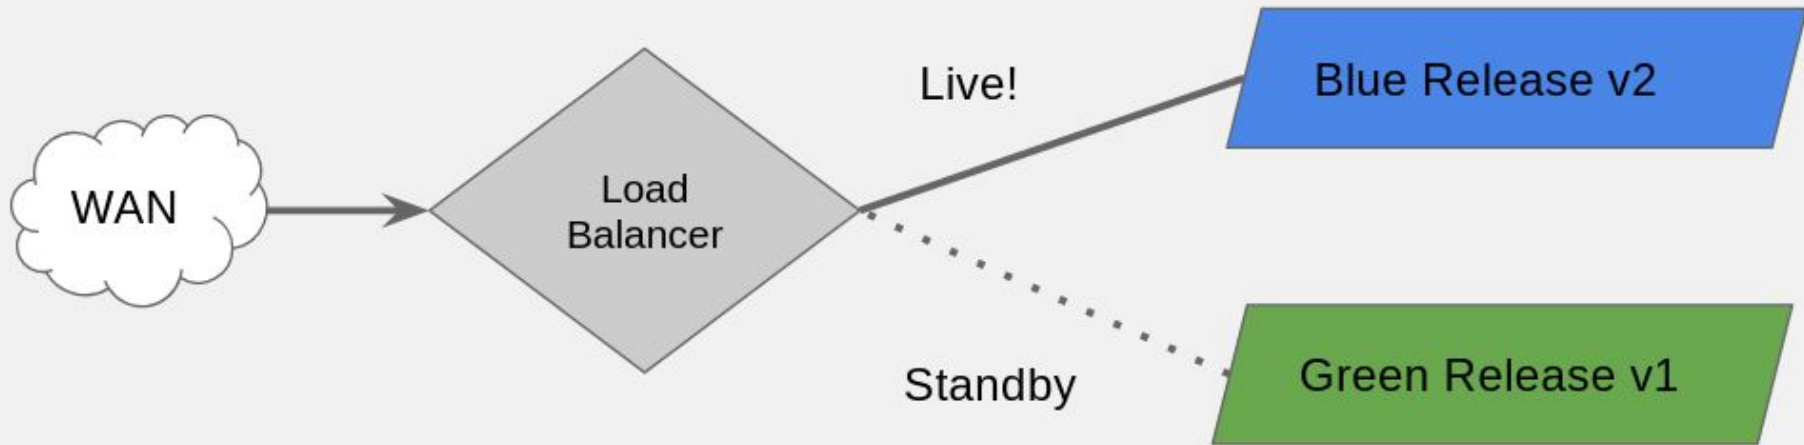


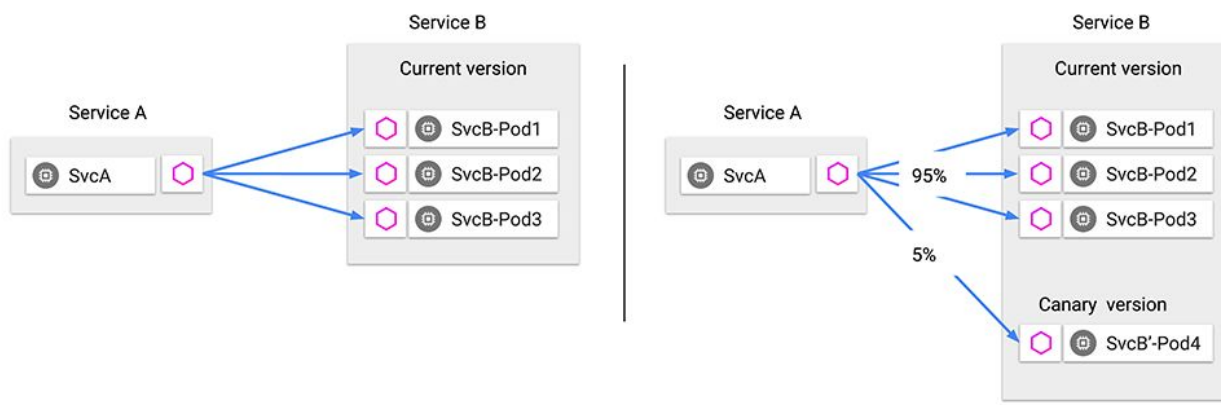
# Microservices'ilities + OpenShift + Istio



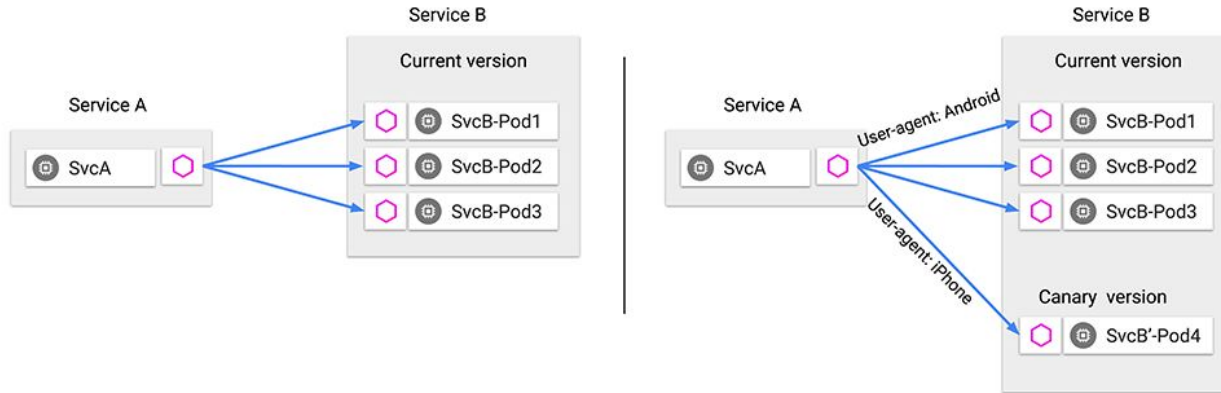
# Blue / Green Deployment

Keep a hot standby ready in case a new release is flawed.

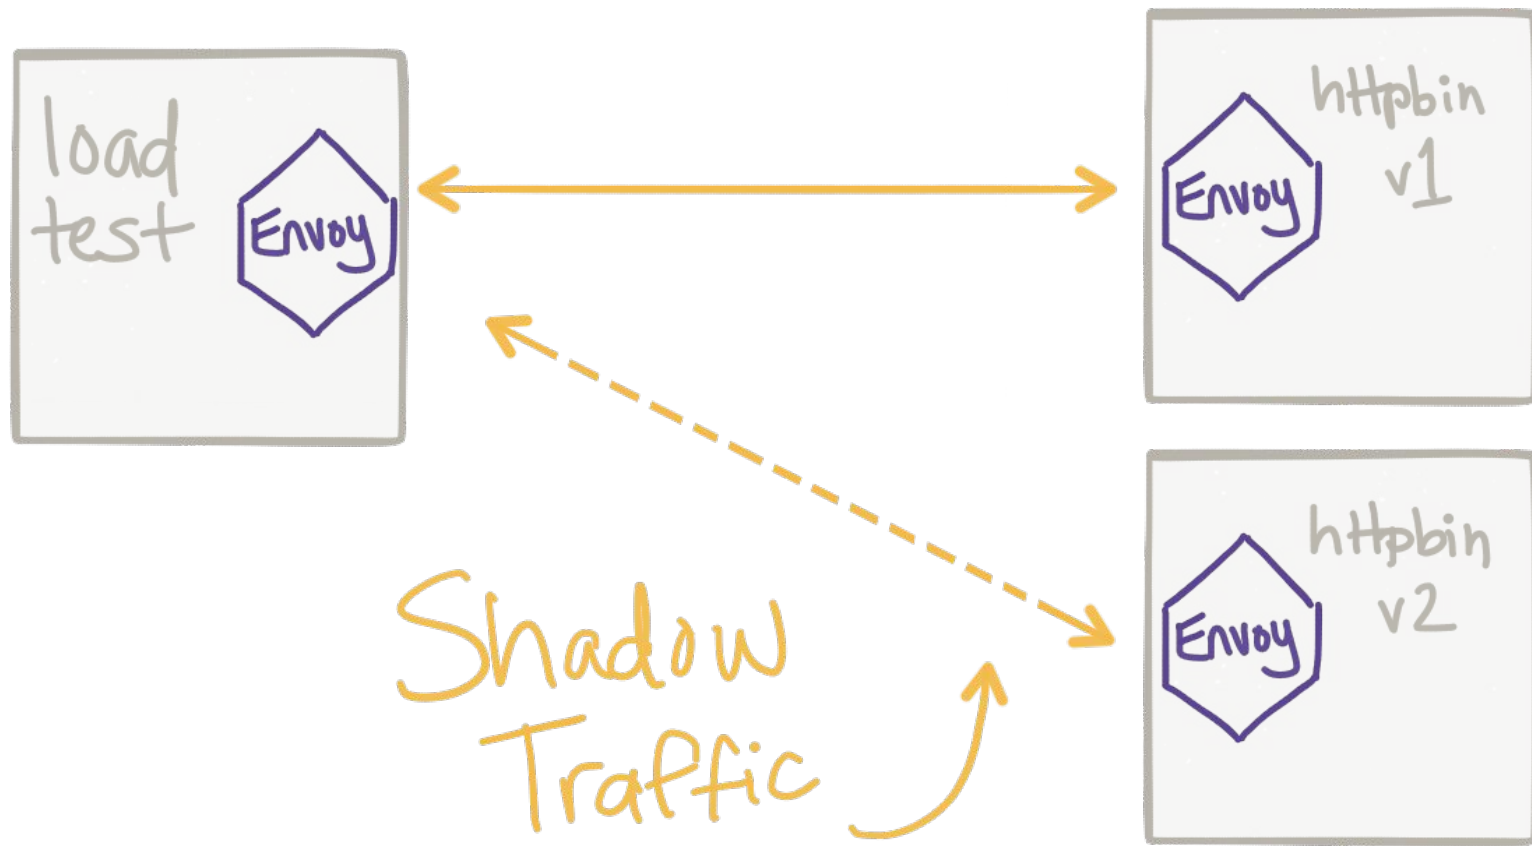




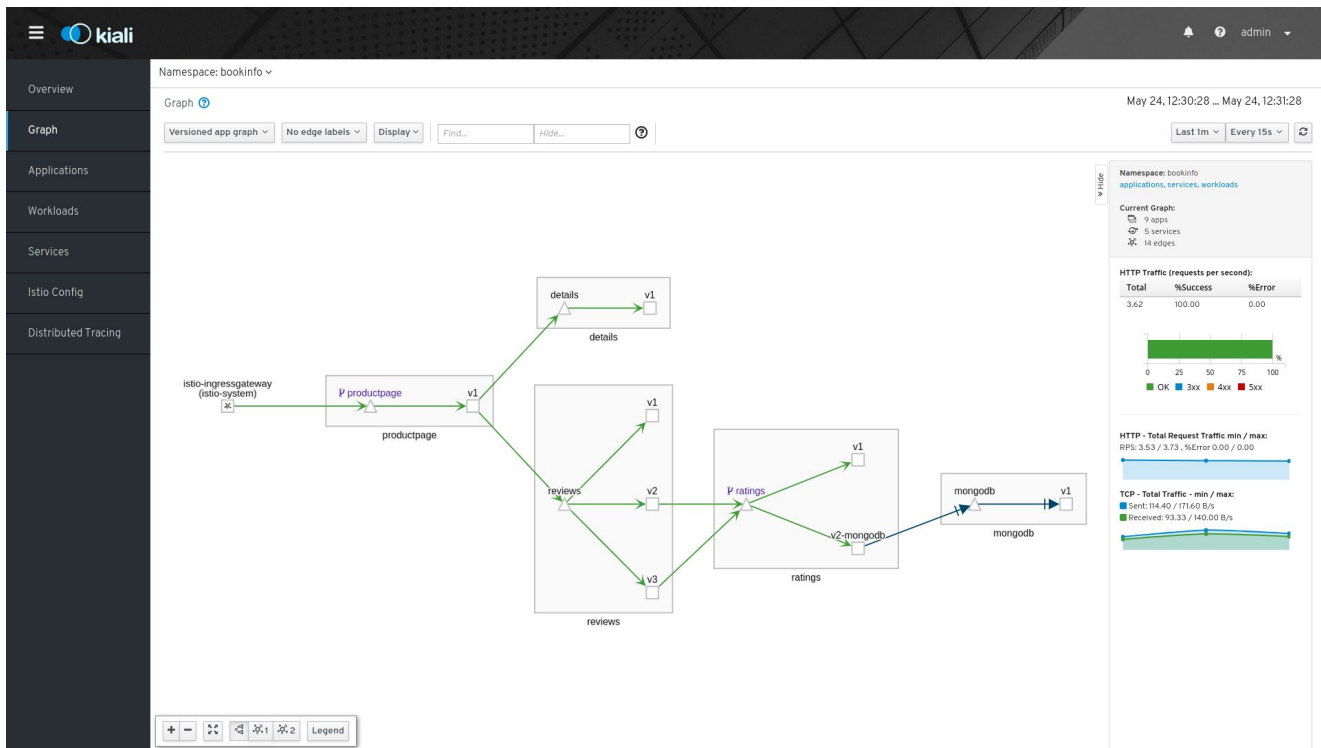
**Traffic splitting decoupled from infrastructure scaling** - proportion of traffic routed to a version is independent of number of instances supporting the version



**Content-based traffic steering** - The content of a request can be used to determine the destination of a request




# Kiali - service mesh observability





# Cloud computing

- IaaS – Infrastructure as a service
  - VMs, servers, storage, network
- PaaS – Platform as a service
  - Execution runtime, database, application server managed Kubernetes, Openshift
- SaaS – Software as a Service
  - Provided applications, CRM, Email, communication



# Cloud-native applications

- Basically microservices
- Designed for cloud deployments
- High requirements on
  - **Low memory utilization**
  - **Low processing requirements**
  - **Fast start-ups**
  - Automation of the app lifecycle
  - CI/CD pipelines

# 12-factor applications

1. **Codebase** – One codebase tracked in revision control, many deploys
2. **Dependencies** – Explicitly declare and isolate dependencies
3. **Config** – Store config in the environment
4. **Backing services** – Treat backing services as attached resources
5. **Build, release, run** – Strictly separate build and run stages

# 12-factor applications

6. **Processes** – Execute the app as one or more stateless processes
7. **Port binding** – Export services via port binding
8. **Concurrency** – Scale out via the process model
9. **Disposability** – Maximize robustness with fast startup and graceful shutdown
10. **Dev/prod parity** – Keep development, staging, and production as similar as possible

# 12-factor applications

- 11. **Logs** – Treat logs as event streams
- 12. **Admin processes** – Run admin/management tasks as one-off processes

<https://12factor.net/>

## Relational DBs



## Metrics / monitoring



## Distributed tracing



JAEGER



## Fault tolerance



HYSTRIX  
DEFEND YOUR APP



Failsafe

## Security



KEYCLOAK

## NoSQL DBs

APACHE  
HBASE



Cassandra



CouchDB  
relax



riak



mongoDB

HYPERTABLE INC



Neo4j




redis


@xstefank

@RedHat


# Demo







# Thank you

-  @xstefank
-  xstefank
- xstefank122@gmail.com

# Resources

- <https://www.zdnet.com/article/to-be-a-microservice-how-smaller-parts-of-bigger-applications-could-remake-it/> originally by Bruce Wong
- <https://medium.com/refraction-tech-everything/how-netflix-works-the-hugely-simplified-complex-stuff-that-happens-every-time-you-hit-play-3a40c9be254b>
- <https://dzone.com/articles/microservices-vs-soa-whats-the-difference>
- <https://martinfowler.com/articles/microservices.html>
- <https://www.docker.com/resources/what-container>
- <https://github.com/kubernetes/kubernetes/blob/master/logo/logo.svg>
- <https://docs.aws.amazon.com/eks/latest/userguide/dashboard-tutorial.html>
- <https://www.slideshare.net/asotobu/service-mesh-patterns>
- [https://access.redhat.com/documentation/en-us/openshift\\_container\\_platform/3.3/html/release\\_notes/release-notes-ocp-3-3-release-notes](https://access.redhat.com/documentation/en-us/openshift_container_platform/3.3/html/release_notes/release-notes-ocp-3-3-release-notes)
- <https://thenewstack.io/history-service-mesh/>
- [https://philcalcado.com/2017/08/03/pattern\\_service\\_mesh.html](https://philcalcado.com/2017/08/03/pattern_service_mesh.html)
- <https://istio.io/docs/concepts/what-is-istio/>
- <http://doughbtv.com/nfvpe/2017/06/05/istio-deploy/>
- <https://blog.aquasec.com/istio-service-mesh-traffic-control>
- <https://blog.christianposta.com/microservices/traffic-shadowing-with-istio-reduce-the-risk-of-code-release/>
- <https://github.com/kiali/kiali>
- <https://blog.openshift.com/what-is-platform-as-a-service-paas/>
- <https://serverless.zone/abstracting-the-back-end-with-faas-e5e80e837362>
- <https://softwareengineeringdaily.com/2016/09/08/relational-databases-with-craig-kerstiens/>
- <https://www.getfilecloud.com/blog/2014/08/leading-nosql-databases-to-consider/>
- <https://www.jaegertracing.io/>
- [https://blog.twitter.com/engineering/en\\_us/a/2012/distributed-systems-tracing-with-zipkin.html](https://blog.twitter.com/engineering/en_us/a/2012/distributed-systems-tracing-with-zipkin.html)
- [https://www.trzcacak.rs/imgm/iTjioIh\\_prometheus-logo-logo-prometheus/](https://www.trzcacak.rs/imgm/iTjioIh_prometheus-logo-logo-prometheus/)
- [https://en.wikipedia.org/wiki/File:Grafana\\_logo.png](https://en.wikipedia.org/wiki/File:Grafana_logo.png)
- <https://design.jboss.org/keycloak/index.htm>
- <https://github.com/Netflix/Hystrix>

# Principles of microservices

# Principles of microservices

- Perspective of business use cases and the solution architecture
- Based on the work of Sam Newman
  - S. Newman, Building Microservices, 1st ed. O'Reilly Media, Inc., 2015
  - S. Newman, “Principles of Microservices,” 2016. [Online]. Available: <http://samnewman.io/talks/principles-of-microservices/>

# 1. Modeled around the business concepts

- Microservices and teams that are maintaining them should correspond to the same business domain
- Microservices are more stable
- Requirements don't change frequently
- Developers are focused on the particular system segment

## 2. Adapting a culture of automation

- With the increasing number of services, their maintenance , administration, and deployment can become unmanageable
- Automation of the is essential
  - Service testing
  - CI/CD
  - Deployment strategy



### 3. Hiding the internal implementation details

- To keep the option of independent development
- Related to bounded-contexts defined by DDD (domain driven design)
  - The context is separated by an explicit interface represented as an API
  - Teams can specify which utilities of the service can be shared and which must be hidden
  - Every request must be processed through this interface

## 4. Decentralizing all things

- Self-sustaining development == services are maintained autonomously
- Decision making delegated to the team maintaining the microservice
- Relevant business logic should be kept in the service and the communication should be as simple as possible

# Conway's law

Any organization that designs a system (defined broadly) will produce a design whose structure is a copy of the organization's communication structure.

# 5. Independent deployments

- Most important principle
- A microservice deployment shouldn't influence the lifespan of any other service
- Various techniques
  - Consumer-driven contracts
  - Test suites for individual parts of the domain
  - CI
  - B/G testing, canary deployments, mirroring

## 6. Customer first

- Service calls must be as simple as possible for customers
- API documentation (Swagger, OpenAPI)
- Service discovery
- Transparency of the call propagation

# 7. Failure isolation

- Fault tolerance to other services failures
- No single point of failure
- Network failures
- Various techniques
  - Fallbacks
  - Retries
  - Timeouts
  - Circuit breakers
  - Bulkheads

# 8. High observability

- Monitoring
  - Individual services
  - System monitoring
- Tracing
  - Synthetic transactions
  - Correlation IDs
- Logging
  - Aggregated logs